

Housing

November 6, 2019

```
In [1]: import pandas
        from matplotlib import pyplot
        import seaborn as sns
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LinearRegression
        from sklearn.linear_model import BayesianRidge
        from scipy import stats
        import numpy
```

```
In [2]: df = pandas.read_csv("datos/housing.csv")
        df.head()
```

```
Out[2]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
In [3]: df.shape
```

```
Out[3]: (20640, 10)
```

```
In [4]: df = df.dropna()
```

```
In [5]: df.shape
```


```
Out[5]: (20433, 10)
```

```
In [6]: df.columns
```

```
Out[6]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
             'total_bedrooms', 'population', 'households', 'median_income',  
             'median_house_value', 'ocean_proximity'],  
            dtype='object')
```

```
In [7]: pyplot.hist(df["median_house_value"], bins = 200)  
        pyplot.xlabel("Price", fontsize=20)  
        pyplot.ylabel("Count", fontsize=20)  
        pyplot.title("Distribution Median House Values")
```

```
Out[7]: Text(0.5, 1.0, 'Distribution Median House Values')
```



```
In [8]: df.corr()["median_house_value"].sort_values()
```

```
Out[8]: latitude           -0.144638  
        longitude          -0.045398  
        population        -0.025300  
        total_bedrooms     0.049686  
        households         0.064894  
        housing_median_age  0.106432  
        total_rooms        0.133294  
        median_income       0.688355  
        median_house_value  1.000000  
        Name: median_house_value, dtype: float64
```

```

In [9]: # One hot encode the variables
dummy_df = pandas.get_dummies(df["ocean_proximity"])

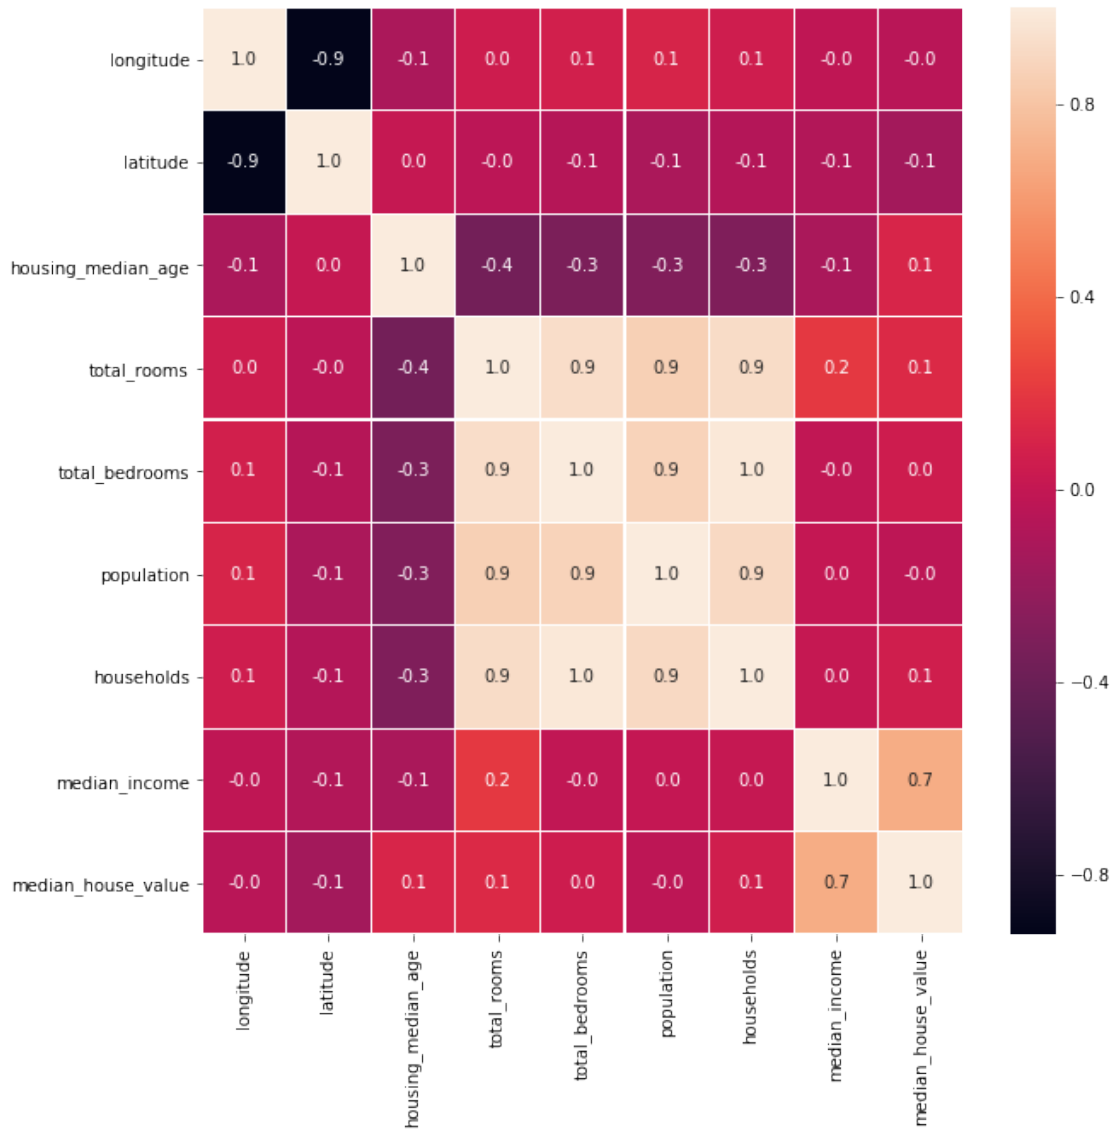
# Put the grade back in the dataframe
dummy_df['median_house_value'] = df['median_house_value']

# Find correlations with grade
dummy_df.corr()['median_house_value'].sort_values()

Out[9]: INLAND          -0.484787
        ISLAND           0.023525
        NEAR OCEAN       0.140378
        NEAR BAY         0.160526
        <1H OCEAN        0.257614
        median_house_value  1.000000
        Name: median_house_value, dtype: float64

In [10]: f, ax = pyplot.subplots(figsize = (10, 10))
sns.heatmap(df.corr(), annot = True, linewidth = 0.3, fmt = ".1f", ax = ax)
pyplot.show()

```



```
In [11]: def format_data(df):
    # Targets are final grade of student
    labels = df['median_house_value']

    # One-Hot Encoding of Categorical Variables
    df = pandas.get_dummies(df)

    # Find correlations with the Grade
    most_correlated = df.corr().abs()['median_house_value'].sort_values(ascending=False)

    # Maintain the top 6 most correlation features with Grade
    most_correlated = most_correlated[:]
```

```

df = df.loc[:, most_correlated.index]

# Split into training/testing sets with 25% split
X_train, X_test, y_train, y_test = train_test_split(df, labels,
                                                    test_size = 0.25,
                                                    random_state=42)

return X_train, X_test, y_train, y_test

```

```

In [12]: X_train, X_test, y_train, y_test = format_data(df)
        X_train.head()

```

```

Out[12]:
      median_house_value  median_income  ocean_proximity_INLAND \
2830                44600.0         1.7109                    1
14951               155000.0         2.4567                    0
8314                450000.0         2.1579                    0
14271                65700.0         1.0531                    0
305                 91200.0         1.8913                    0

      ocean_proximity_<1H OCEAN  ocean_proximity_NEAR BAY  latitude \
2830                        0                0          35.40
14951                       1                0          32.71
8314                        0                0          33.35
14271                       0                0          32.70
305                         0                1          37.76

      ocean_proximity_NEAR OCEAN  total_rooms  housing_median_age \
2830                        0        8739.0          11.0
14951                       0        2413.0          18.0
8314                        0        1675.0          27.0
14271                       1         818.0          38.0
305                         0        2018.0          43.0

      households  total_bedrooms  longitude  population \
2830        1919.0         2190.0    -119.01        4781.0
14951         551.0         533.0    -116.96        1129.0
8314         331.0         521.0    -118.32         744.0
14271         231.0         217.0    -117.12         953.0
305          367.0         408.0    -122.18        1111.0

      ocean_proximity_ISLAND
2830                        0
14951                       0
8314                        1
14271                       0
305                         0

```

```

In [13]: print(X_train.shape)

```

```

        print(X_test.shape)

(15324, 14)
(5109, 14)

In [14]: def evaluate(X_train, X_test, y_train, y_test):
    # Names of models
    model_name_list = ['Linear Regression', "Bayesian Ridge"]

    X_train = X_train.drop(columns='median_house_value')
    X_test = X_test.drop(columns='median_house_value')

    # Instantiate the models
    model1 = LinearRegression()
    model2 = BayesianRidge(compute_score=True)

    # Dataframe for results
    results = pandas.DataFrame(columns=['mae', 'rmse'], index = model_name_list)

    # Train and predict with each model
    for i, model in enumerate([model1, model2]):
        model.fit(X_train, y_train)
        predictions = model.predict(X_test)
        coeff = model.coef_
        score = model.score(X_test, y_test)

        print("{} Coefficients: ".format(model_name_list[i]), coeff)
        print("\n")
        print("{} Score: ".format(model_name_list[i]), score)
        print("\n")
        # Metrics
        mae = numpy.mean(abs(predictions - y_test))
        rmse = numpy.sqrt(numpy.mean((predictions - y_test) ** 2))

        # Insert results into the dataframe
        model_name = model_name_list[i]
        results.loc[model_name, :] = [mae, rmse]

    # Median Value Baseline Metrics
    baseline = numpy.mean(y_train)
    baseline_mae = numpy.mean(abs(baseline - y_test))
    baseline_rmse = numpy.sqrt(numpy.mean((baseline - y_test) ** 2))

    results.loc['Baseline', :] = [baseline_mae, baseline_rmse]

    return results

In [15]: results = evaluate(X_train, X_test, y_train, y_test)

```

```
Linear Regression Coefficients: [ 3.92330777e+04 -7.36819088e+04 -3.40321348e+04 -4.02817945e+04  
-2.55205753e+04 -3.12916376e+04 -6.36691750e+00 1.08813173e+03  
4.35107705e+01 1.02597469e+02 -2.68384898e+04 -3.62122333e+01  
1.79287476e+05]
```

```
Linear Regression Score: 0.6547516410329023
```

```
Bayesian Ridge Coefficients: [ 3.92244863e+04 -5.00655360e+04 -1.05743626e+04 -1.67689497e+04  
-2.56126711e+04 -7.85132337e+03 -6.36482591e+00 1.09028421e+03  
4.30517739e+01 1.03031260e+02 -2.69134708e+04 -3.62257024e+01  
8.52601716e+04]
```

```
Bayesian Ridge Score: 0.6553914482946943
```

```
In [16]: print(results)
```

	mae	rmse
Linear Regression	49958.5	68343.3
Bayesian Ridge	49921.4	68279.9
Baseline	92461.7	116314

```
In [17]: X_test.iloc[50]
```

```
Out[17]: median_house_value    177500.0000  
         median_income          2.4827  
         ocean_proximity_INLAND    0.0000  
         ocean_proximity_<1H OCEAN  1.0000  
         ocean_proximity_NEAR BAY    0.0000  
         latitude                38.3600  
         ocean_proximity_NEAR OCEAN  0.0000  
         total_rooms              5496.0000  
         housing_median_age        6.0000  
         households               1189.0000  
         total_bedrooms            1374.0000  
         longitude                -122.6900  
         population               2502.0000  
         ocean_proximity_ISLAND    0.0000  
         Name: 19157, dtype: float64
```

```
In [18]: model1 = LinearRegression()  
         model2 = BayesianRidge(compute_score=True)
```

```
# Train and predict with each model
for i, model in enumerate([model1, model2]):
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    print(predictions[50])
```

```
177499.99999999997
```

```
177500.00000951337
```

```
In [ ]:
```

```
In [ ]:
```